

CS323 Operating Systems Memory Management III

Yuanyuan Zhou
Lecture 15
2/24/2003

Content of this lecture

- Administrative announcements
- Page Table and TLB
- Multi-level Paging
- Inverted Page Table
- Protection
- Summary

2

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Administrative

- Quiz2 starts
- MP2

3

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Review

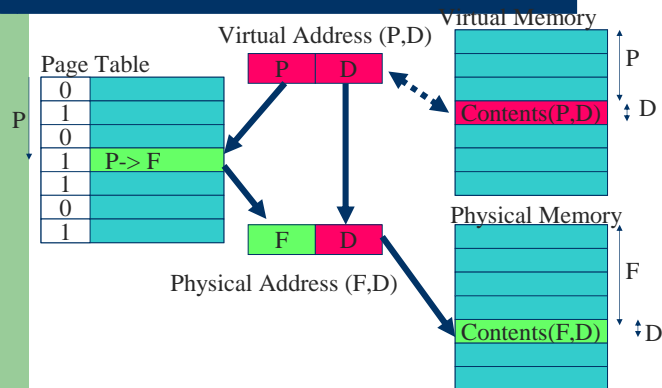
- Storage management
 - Bitmap or link list
 - Compaction
 - Best fit, quick fit, first fit, next fit, worst fit
- Virtual Memory
- Paging
 - Page fault

4

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Page Mapping Hardware (page hit)

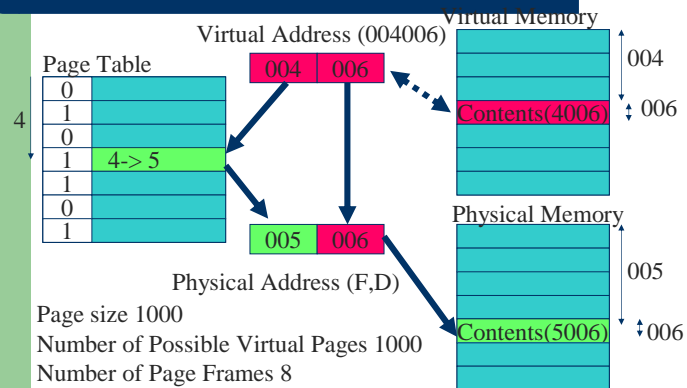


5

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Page Mapping Hardware



6

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Page Fault

- Access a virtual page that is not mapped into any physical page
 - A fault is triggered by hardware
- Page fault handler (by VM Software, a part of OS)
 - Find if there is any free physical page available
 - If no, evict some resident page to disk (swapping space)
 - Allocate a free physical page
 - Load the faulted virtual page to the prepared physical page
 - Modify the page table

7

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Paging Issues

- Page size is 2^n
 - usually 512, 1k, 2k, 4k, or 8k
 - E.g. 32 bit VM address may have 2^{20} (1MB) pages with 4k (2^{12}) bytes per page
- Page table:
 - 2^{20} page entries take 2^{22} bytes (4MB)
 - page frames must map into real memory
 - Page Table base register must be changed for context switch
- NO External fragmentation, Internal fragmentation on last page ONLY

8

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Virtual-To-Physical Lookups


- Programs only know virtual addresses
 - The page table can be extremely large
- Each virtual address must be translated
 - May involve walking hierarchical page table
 - Page table stored in memory
 - So, each program memory access requires several actual memory accesses
- Solution: cache “active” part of page table
 - TLB also called “associative memory”

9

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Translation Look-aside Buffer (TLB)



10

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

TLB Function


- If a virtual address is presented to MMU, the hardware checks TLB by comparing all entries simultaneously (in parallel).
- If match is valid, the page is taken from TLB without going through page table.
- If match is not valid
 - MMU detects miss and does an ordinary page table lookup.
 - It then evicts one page out of TLB and replaces it with the new entry, so that next time that page is found in TLB.

11

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Page Mapping Hardware




12

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Page Mapping Example




13

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Page Mapping Example: next reference



14

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Bits in a TLB Entry

- Common (necessary) bits
 - Virtual page number: match with the virtual address
 - Physical page number: translated address
 - Valid
 - Access bits: kernel and user (nil, read, write)
- Optional (useful) bits
 - Process tag
 - Reference
 - Modify
 - Cacheable

15

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Paging Implementation Issues

- TLB can be implemented using
 - Associative registers
 - Look-aside memory
 - Content-addressable memory
- TLB hit ratio (Page address cache hit ratio)
 - Percentage of time page found in associative memory

16

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Hardware-Controlled TLB

- On a TLB miss (different from page fault)
 - Hardware loads the PTE into the TLB
 - Need to write back if there is no free entry
 - Generate a fault if the page containing the PTE is invalid
 - VM software performs fault handling
 - Restart the CPU
- On a TLB hit, hardware checks the valid bit
 - If valid, pointer to page frame in memory
 - If invalid, the hardware generates a page fault
 - Perform page fault handling
 - Restart the faulting instruction

17

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Software-Controlled TLB

- On a miss in TLB, VM software
 - Write back if there is no free entry
 - Check if the page containing the PTE is in memory
 - If no, perform page fault handling
 - Load the PTE into the TLB
 - Restart the faulting instruction
- On a hit in TLB, the hardware checks valid bit
 - If valid, pointer to page frame in memory
 - If invalid, the hardware generates a page fault
 - Perform page fault handling
 - Restart the faulting instruction
- Example: RISC including SPARC, Alpha, MIPS, HP PA

18

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Hardware vs. Software Controlled

- Hardware approach
 - Efficient
 - Inflexible
 - Need more space for page table
- Software approach
 - Flexible
 - Software can do mappings by hashing
 - PP# → (Pid, VP#)
 - (Pid, VP#) → PP#
 - Can deal with large virtual address space

19

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Group Discussion

- Similarity between Cache and TLB
- Difference between Cache and TLB

20

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Cache vs. TLBs

- Similarities
 - Both cache a portion of memory
 - Both write back on a miss
- Combine L1 cache with TLB
 - Virtually addressed cache
 - Why wouldn't everyone use virtually addressed caches?
- Differences
 - Associativity
 - TLB is usually fully set-associative
 - Cache can be direct-mapped
 - Consistency
 - TLB does not deal with consistency with memory
 - TLB can be controlled by software

21

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Issues

- What TLB entry to be replaced?
 - Random
 - Pseudo Least Recently Used (LRU)
- What happens on a context switch?
 - Process tag: change TLB registers and process register
 - No process tag: Invalidate the entire TLB contents
- What happens when changing a page table entry?
 - Change the entry in memory
 - Invalidate the TLB entry

22

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Effective Access Time

- TLB lookup time = ϵ time unit
- Memory cycle -- m microsecond
- TLB Hit ratio -- α
- Effective access time
 - $Eat = (1m + \epsilon)\alpha + (2m + \epsilon)(1 - \alpha)$
 - $Eat = 2m + \epsilon - \alpha$

23

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Multilevel Page Tables


- Since the page table can be very large, one solution is to page the page table
- Divide the page number into
 - An index into a page table of second level page tables
 - A page within a second level page table
- Advantage
 - No need to keeping all the page tables in memory all the time
 - Only recently accessed memory's mapping need to be kept in memory, the rest can be fetched on demand

24

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Multiple-Level Page Tables



What does this buy us? Sparse address spaces and easier paging


25

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Example Addressing on a Multilevel Page Table System

- A logical address (on 32 bit machine with 4k page size) is divided into
 - A page number consisting of 20 bits
 - A page offset consisting of 12 bits
- Divide the page number into
 - A 10-bit page number
 - A 10-bit page offset



26

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Multilevel Paging and Performance


- Since each level is stored as a separate table in memory, converting a logical address to a physical one in a four-level paging may take five memory accesses. Why?

27

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Inverted Page Tables




- Main idea
 - One PTE for each physical page frame
 - Hash (Vpage, pid) to Ppage#
 - Trade off space for time
- Pros
 - Small page table for large address space
- Cons
 - Lookup is difficult
 - Overhead of managing hash chains, etc

28

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Inverted Page Table




29

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Inverted Page Table Implementation

- TLB is same as before
- TLB miss is handled by software
- In-memory page table is managed using a hash table
 - number of entries = number of physical frames
 - Not found: page fault



30

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Sharing Pages


- Code and data can be shared among processes
 - mapping them into pages with common page frame mappings
- Code and data must be position independent if VM mappings for the shared data are different
- Code and data cannot store VM addresses for functions and variables not in shared pages
- Shared code and data usually are not position independent resulting in certain regions of memory being reserved for shared libraries and the operating system

31

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Shared Pages

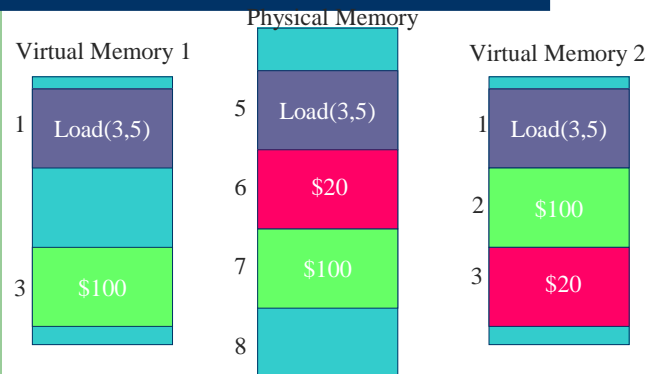


32

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Incorrect Sharing



33

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Protection

- Can add read, write, execute protection bits to page table to protect memory
- Check is done by hardware during access
- shared memory location
 - different protections from different processes
 - Solution:
 - associate protection lock with page frame. Each process has its own key. If the key fits the lock, the process may access the page frame

34

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Protection


- Typically many different keys can fit a lock using a priority numbering scheme
- (E.G. Key 3 fits all locks 3, 7, 15)
- Key 4 fits 5,6,7,12,13,14,15.)

35

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Keys and Locks



36

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou

Reminder

- MP2

37

2/24/2003

CS 323 - Operating Systems,
Yuanyuan Zhou