

University of Illinois at Urbana-Champaign
Department of Computer Science

Midterm Examination 1

CS 323 Operating System Design
Spring, 2003

7-8pm Monday, March 10th
DCL 1320 & MSEB 100

Print your name and ID number neatly in the space provided below; print your name at the upper right corner of every page.

Name:
Net ID:

This is a **closed book, closed notes, No calculator** exam.

Write your answers **CLEARLY**. The answer is wrong if the instructor and the TAs cannot read it.

Do all parts of all four problems in this booklet. This booklet should include this title page, 7 additional pages and 2 empty pages. Do your work inside this booklet, using the empty pages and backs of pages if needed. Problems are of various difficulty degree, hence if you don't know the answer immediately, progress to the next problem and come back to the unsolved problem later.

Problem	Score	Grader
1		
2		
3		
4		
Total		

1. Problem – Process Scheduling (25 Points)

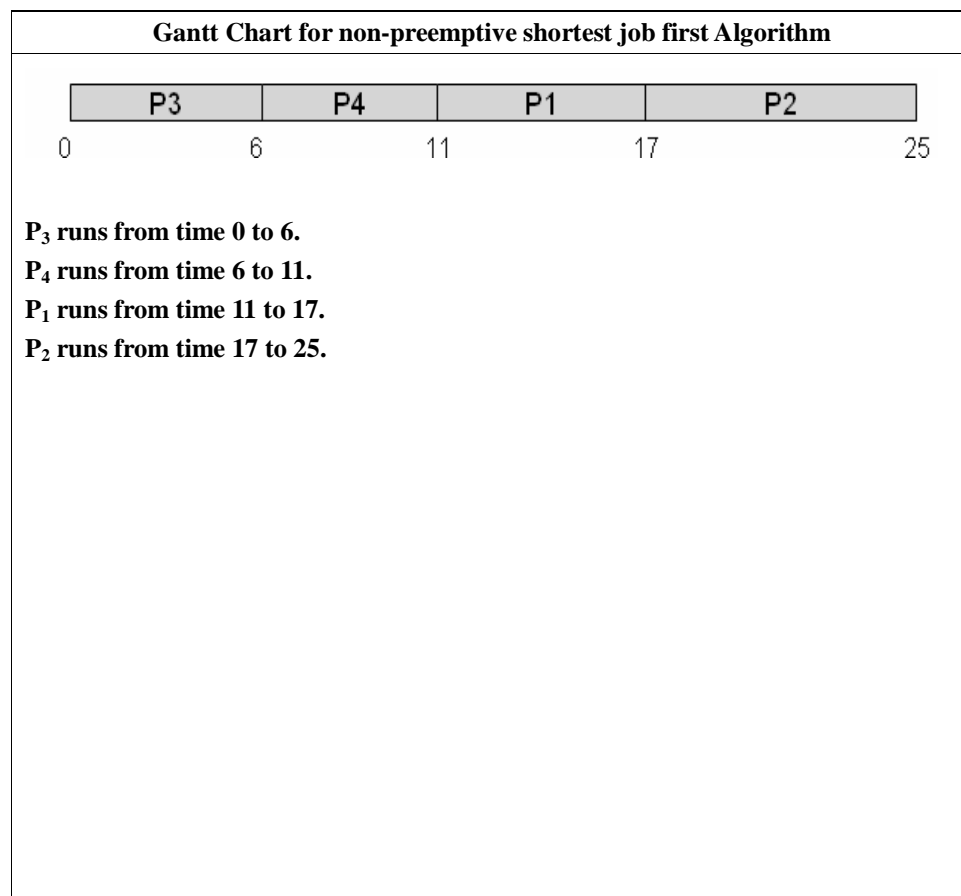
Suppose we have 4 processes arrive for execution. The process burst time (duration), as well as the arrival time and priority of each process is indicated in the following table. Time is in milliseconds (ms).


Process	Arrival Time	Burst Time (duration)	Priority
P ₁	6	6	4
P ₂	3	8	2
P ₃	0	6	3
P ₄	1	5	1

Table 1: The processes with their arrival time, process behavior and priority

1. (15 Points) Fill out the three Gantt charts below illustrating the scheduled execution of these processes using
 - a) non-preemptive shortest job first scheduling algorithm
 - b) Round Robin (quantum = 4ms)
 - c) Preemptive priority (a smaller priority number implies a higher priority).

Note: CPU cannot stay idle if there are still unfinished processes.





2. Problem-Synchronization (20 points)

1. (10 points) Show how a binary semaphore can be implemented using atomic SWAP instruction.

The following is the definition of swap instruction:

```
void Swap (char* x,* y) \\ All done atomically
{
    char temp = *x;
    *x = *y;
    *y = temp
}
```

```
typedef int mutex;
Mutex m = 1; // initially mutex is free
Semaphore::P(Mutex *m) {
{
    char temp = 0;

    swap(m, &temp );
    while ( temp == 0 ) {
        swap(m, &temp );
    }
}

Semaphore::V(mutex *m)
{
    char temp = 1;

    swap(m, &temp );
}
```

2. (10points) Consider the following producer and consumer code. Is the code correct? If yes, explain what is happening in each line of the code. If no, specify: what is wrong, which line is wrong and why it is wrong. The number at the very beginning of each line is the line number.

```

1. typedef int semaphore;
2. semaphore mutex=1;
3. semaphore empty = N;
4. semaphore full = 0;
5.
6. void producer(void)
7. { int item;
8.
9.   while (TRUE) {
10.      item = produce_item ();
11.      down(&mutex);
12.      insert_item(item);
13.      down(&empty);
14.      up(&full);
15.   }
16. }
17. void consumer(void)
18. { int item;
19.
20.   while (TRUE) {
21.      down(&mutex);
22.      item= remove_item();
23.      down(&full);
24.      up(&empty);
25.      consume_item(item);
26.   }
27. }
```

There are four problems in the code. These are identified, together with appropriate fixes, below.

Line 10/11: Before the *down* on *mutex*, the producer should also *down* the *empty* semaphore. This will block the producer if the number of empty slots is 0 (that is, there is no space left in the shared buffer). Note that this must be done before the *down* on *mutex* to avoid the potential for deadlock.

Line 13. Doing a *down* on the *empty* semaphore at this point makes no sense. The producer should instead release the mutual exclusion it holds by doing an *up* on *mutex*.

Line 20/21. This is an analogous problem to that between lines 10 and 11. The consumer must, before doing *down* on the *mutex* semaphore, also *down* the *full* semaphore. This will block it if the number of full slots is 0 (all slots are empty, in other words).

Line 23. The *down* on the *full* semaphore is not correct. Instead, an *up* on the *mutex* semaphore should happen here, indicating the consumer is leaving the critical section.

3 Problem-Deadlock (30points)

Consider two resource types A and B. Furthermore, consider the snapshot of a system in the following table. Allocation matrix (column) lists the number of resources allocated for each process; while the Max. Request matrix (column) lists the maximum number resources each process request.

	Allocation		Max. Request		Available		Need	
	A	B	A	B	A	B	A	B
P ₀	2	0	2	5	3	Y	0	5
P ₁	5	2	10	2			5	0
P ₂	1	2	5	4			4	2
P ₃	1	1	4	1			3	0
P ₄	0	4	5	9			5	5

- (5 points) Determine the Need Matrix (column) for resource types A and B and fill them in then above table.

See Above

- (10points) Suppose Y=6, is it a safe state? If it is safe, provide a safe sequence for executing all processes (In order to get a full credit, show your work how you find the safe sequence). If it is not safe, explain it.

The state is a safe one. There are multiple correct sequences. One possibility is shown below. The columns show how many of A and B is available after that process successfully executes. Essentially the constraints are that either P₀ or P₃ must run first. If P₀ runs first, the rest of the processes can run in any order. If P₃ runs first, P₀ or P₂ must run next.

Available		
	A	B
	3	6
P ₀	5	6
P ₁	10	8
P ₂	11	10
P ₃	12	11
P ₄	12	15

3. (15 points) What is the smallest value of Y for which this is a safe state? Provide a safe sequence for executing all processes for the situation with the smallest value of Y for which this is a safe state. For your convenience, the table in the previous page is **copied** below:

	Allocation		Max. Request		Available		Need	
	A	B	A	B	A	B	A	B
P ₀	2	0	2	5	3	Y	0	5
P ₁	5	2	10	2			5	0
P ₂	1	2	5	4			4	2
P ₃	1	1	4	1			3	0
P ₄	0	4	5	9			5	5

The need matrix is as above. We compute minimum Y by first assuming Y=0 and finding a process that we can execute. If one is not available, we pick the process with the minimum B requirement. We repeat the algorithm until all processes have executed. (Note: multiple correct sequences may exist).

	A	B	B Need
	3	Y	
P ₃	4	Y+1	0
P ₂	5	Y+3	2
P ₁	10	Y+5	0
P ₀	12	Y+5	5
P ₄	12	Y+9	5

We can now compute Y in the following manner. We look at the first row and see that to execute P₃ we need 0 B, thus Y=0 (from the row above). Next, to execute P₂, Y+1=2, or Y=1. We proceed like this and find that Y must be at least 1 to allow this sequence to execute.


The above is a safe sequence

4 Problem – Memory Management (25 Points)

In most systems, a process can allocate and de-allocate memory from its heap using *malloc()* and *free()*. A heap is just a chunk of memory. Suppose the heap is managed with linked list. Each node in the list is either allocated or free. **And the list kept sorted by address in increasing order.** When *malloc()* is called, the list is searched for a free segment that is big enough (depending on the allocation algorithm), that segment is **divided** into an allocated segment (at the beginning) and a free segment. When *free()* is called, the corresponding segment should **merge** with its neighboring segments, if they are also free.

Now suppose a process has a heap of 13KB, which is initially empty. During its execution, the process issues the following memory allocate/de-allocate calls (pi is just a pointer):

```
p1 = malloc(3KB);
p2 = malloc(4KB);
p3 = malloc(1KB);
free(p2);
p4 = malloc(2KB);
free(p1);
p5 = malloc(2KB);
```

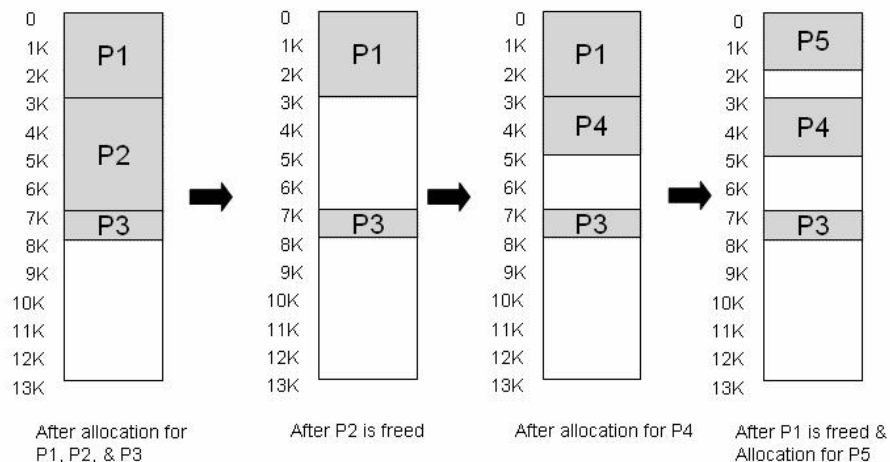


Show what the heap is like after the above calls, using


- (1) first-fit (5 points),
- (2) best-fit (10 points), and
- (3) worst-fit (10 points)

algorithms respectively. For simplicity, assume the memory begins at address 0, and ignore the memory used by the linked list itself.

(1) First-fit



(2) Best-fit



(3) Worst-fit

